

```
In [6]: #!pip install openai langchain faiss-cpu sentence-transformers  
#!pip install -U langchain-community
```

```
In [ ]: # Step 1: Import Required Libraries for LLM + Document Retrieval Workflow  
import os  
import torch  
from transformers import pipeline  
from langchain.document_loaders import TextLoader  
from langchain.text_splitter import RecursiveCharacterTextSplitter  
from langchain.embeddings import HuggingFaceEmbeddings  
from langchain.vectorstores import FAISS  
from langchain.llms import HuggingFacePipeline  
from langchain.chains import RetrievalQA
```

LangChain is a framework for building applications with language models (LLMs).

- It simplifies the process of integrating LLMs with external data sources (like documents, databases, or APIs)
- It allows you to create intelligent apps such as chatbots, retrieval-augmented generation (RAG) systems, and more.

LangChain acts as the "glue" that connects:

- Your documents (TextLoader)-----> Loads plain text documents into LangChain.
- Your text processing (TextSplitter)-----> Splits long text into manageable chunks.
- Your embeddings (HuggingFaceEmbeddings)-----> Generates embeddings for document chunks.
- Your retrieval system (FAISS)-----> Stores and retrieves embeddings efficiently using similarity search. It helps you find the most similar text chunks from your documents when someone asks a question.
- Your LLM (HuggingFacePipeline)-----> Wraps a Hugging Face model as an LLM.
- And your question-answering logic (RetrievalQA)-----> Connects retrieval and LLM to answer questions based on documents.

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```

In [11]: # Step 2: Load the document. A text loader is used to load the document as raw text into memory
loader = TextLoader("llm_notes.txt", encoding="utf-8")
documents = loader.load()

# Step 3: Split the document. Each chunk has 300 characters, and 50 characters overlap with the previous chunk
text_splitter = RecursiveCharacterTextSplitter(chunk_size=300, chunk_overlap=50)
docs = text_splitter.split_documents(documents[:5]) # limit size to reduce memory

# Step 4: Embed using Hugging Face sentence transformer
embeddings = HuggingFaceEmbeddings(model_name="all-MiniLM-L6-v2")
db = FAISS.from_documents(docs, embeddings)

# Step 5: Load flan-t5-small on CPU (safest config). Hugging Face's pipeline is wrapped into a PyTorch pipeline
device = torch.device("cpu")

qa_pipeline = pipeline(
    "text2text-generation",
    model="google/flan-t5-small",
    max_length=256,
    device=device,
    do_sample=False
)

llm = HuggingFacePipeline(pipeline=qa_pipeline)

# Step 6: Build the Retrieval QA chain. First retrieves top 3 relevant text chunks from FAISS,
# It enhances the LLM's ability to answer questions by grounding it in specific documents.
retriever = db.as_retriever(search_kwargs={"k": 3})
qa_chain = RetrievalQA.from_chain_type(
    llm=llm,
    retriever=retriever,
    return_source_documents=True
)

# Step 7: Ask a question
query = "What is tokenization in LLMs and why is it important?"
result = qa_chain(query)

# Step 8: Show results
print("\n Answer:")
print(result["result"])

print("\n Source Documents:")
for i, doc in enumerate(result["source_documents"], 1):
    print(f"\n--- Source {i} ---")
    print(doc.page_content)

```

Device set to use cpu

Answer:

Tokenization is the process of splitting text into smaller units called tokens.

Source Documents:

--- Source 1 ---

Tokenization is the process of splitting text into smaller units called tokens.

Large language models typically use subword tokenization for efficiency.

Example: 'unbelievable' → ['un', 'believ', 'able'].

In []: