In [3]:
```
#!pip install transformers
```

In [5]:
```
from transformers import BertTokenizer, BertForSequenceClassification
import torch

# Load the pre-trained BERT tokenizer and model
tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
model = BertForSequenceClassification.from_pretrained("bert-base-uncased", num_labels=2)  # Bi

# Define a text for classification
text = "I really enjoyed the movie. The acting was fantastic!"

# Tokenize the text and convert to input features
inputs = tokenizer(text, return_tensors="pt")
input_ids = inputs["input_ids"]
attention_mask = inputs["attention_mask"]

# Perform inference
with torch.no_grad():
    outputs = model(input_ids, attention_mask=attention_mask).logits

# Get predicted label
predicted_label = torch.argmax(outputs, dim=1).item()

# Define label names
label_names = ["Negative", "Positive"]

# Print result
print("Text:", text)
print("Predicted Label:", label_names[predicted_label])
```

```
Some weights of BertForSequenceClassification were not initialized from the model checkpoint
at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predicti
ons and inference.

Text: I really enjoyed the movie. The acting was fantastic!
Predicted Label: Positive
```

In [3]: